

MATHEMATICAL MODELS FOR NATURAL GAS FORECASTING

STEVEN R. VITULLO, RONALD H. BROWN¹,
GEORGE F. CORLISS AND BRIAN M. MARX

ABSTRACT. It is vital for natural gas Local Distribution Companies (LDCs) to forecast their customers' natural gas demand accurately. A significant error on a single very cold day can cost the customers of the LDC millions of dollars. This paper looks at the financial implication of forecasting natural gas, the nature of natural gas forecasting, the factors that impact natural gas consumption, and describes a survey of mathematical techniques and practices used to model natural gas demand. Many of the techniques used in this paper currently are implemented in a software GasDayTM, which is currently used by 24 LDCs throughout the United States, forecasting about 20% of the total U.S. residential, commercial, and industrial consumption. Results of GasDay'sTM forecasting performance also is presented.

1 Introduction A natural gas Local Distribution Company (LDC) faces many challenges in the business of supplying gas to its customers. The gas supply system of an LDC consists of gate stations, compressors, gas storage, and customers. The LDC must operate these systems to assure delivery of gas in adequate volumes at required pressures under all circumstances. For efficient, economical, and safe operation, the daily gas demanded by the customers must be known in advance with a relatively high degree of accuracy. Similar models are used to forecast hourly demands and also monthly and longer term demands. This paper discusses methods to predict aggregate daily demand of the customers of an LDC.

The customer base of an LDC consists of many individual customers, each with unique demand characteristics. Customers use gas for space heating, known as heating load, for heating water, drying, cooking and baking, and other processes, known as base load, and for electric power

¹Director, GasDayTM Project.

Keywords: Neural networks, regression, utility forecasting.

Copyright ©Applied Mathematics Institute, University of Alberta.

generation. Heating load is dependent on weather (most importantly temperature) factors that affect consumption. Meanwhile, base load accounts for other factors that are not weather dependent and tend to be constant, although it may change over time with growth in customer base. Heating load is challenging to forecast as it requires forecasts of weather factors, and base load is difficult to forecast because it requires knowledge of customer behavior. The customer base generally is divided into four categories: residential, commercial, industrial, and electric power generation [1]. In this paper, we discuss just residential, commercial, and industrial demand. The demand characteristics of these three categories differ significantly. The residential customer demands are typically temperature sensitive with increasing consumption on weekends. A commonly used measurement for natural gas energy consumption is a decatherm, equal to one million British thermal units [1]. A typical gas-heated Wisconsin residence consumes approximately one decatherm of gas on a cold winter day. Commercial customers tend to be temperature sensitive and decrease their gas use on weekends. Industrial customers tend to be less temperature sensitive but also have decreasing consumption on weekends. Additionally, customers are subdivided into two service contract groups. A firm customer has a service contract which anticipates no service interruptions, and an interruptible customer has a service contract that allows the LDC to interrupt service during peak demand times. We do not discuss or model electric power generators because electric power generated at one end of the country can be transported to the other end of the country for end use. This makes forecasting natural gas demand for electric power generation a fundamentally different forecasting problem outside the scope of this paper.

LDCs are required by state utility commissions to supply uninterrupted gas service to their firm customers in a cost-effective manner during a peak day, the day on which maximum gas loads are experienced. Since a large portion of natural gas is consumed for space heating, natural gas consumption in many operational areas is heavily weather-dependent [46]. Thus, the peak load day is likely to occur during the coldest weather conditions. For example, a large natural gas utility may have a heat-dependent load of approximately 10,000 decatherms per degree Fahrenheit. This means that approximately 10,000 additional decatherms of natural gas are consumed (for heating purposes) for each degree Fahrenheit colder it gets. The contract price of natural gas has varied in recent years from approximately \$4.00 to approximately \$15.00 per decatherm, whereas the spot market price of gas on a high demand

day can be 10 times the contract price [17]. Thus, in this example, \$400,000 to \$1,500,000 of additional cost is introduced for each degree Fahrenheit the LDC overestimated the temperature during extreme cold weather conditions, assuming the gas was purchased on the spot market at 10 times the cost of gas purchased under contract. The large cost the LDC incurs for buying gas on the spot market is passed directly to the end customer. Accurate forecasting of air conditioning loads and of non-temperature dependent gas demand (commonly referred to as baseload gas) during warm weather conditions is equally important.

Historically, many methods have been used to predict daily demand [29, 32, 33]. Gas controllers have used methods such as looking at use patterns on similar historical days and scatter plots of use versus temperature. Often these methods are applied successfully only by experts with years of experience at the LDC.

Along with deregulation of gas prices came the need to forecast customer demand for natural gas more accurately than before. Although many of the larger LDCs have the ability to store or withdrawal gas to cover their forecast error, the majority of LDCs do not have storage capabilities, making their forecast accuracy critical. Many LDC's have developed mathematical formulas to predict gas demand with varying degrees of success. These models are developed using historical demand data and other historical data and information, such as weather conditions and day of the week.

2 Mathematical models to forecast daily demand The most common mathematical modeling techniques used to forecast daily demand are multiple linear regression and artificial neural networks. This section briefly presents these two methods used by GasDayTM, a forecasting software application licensed to 24 LDCs in the US. Section 3 and 4 present the factors used by GasDayTM that affect daily natural gas demand and data quality, respectively. These sections focus on the things we considered when building GasDayTM. Section 5 presents an analysis of the performance of GasDay'sTM forecasts.

2.1 Multiple linear regression Multiple Linear Regression (MLR) [15, 18] is one of the most commonly used methods for prediction models, and it has been applied to utility forecasting [19]. Suppose for N days ($1 \leq k \leq N$), we have customer demand S_i and M independent factors, $x_{k,j}$, for $1 \leq k \leq N$ and $1 \leq j \leq M$ we think may affect S_i . The

multiple linear regression model estimates

$$S_k \approx \hat{S}_k = \beta_0 + \sum_{j=1}^m \beta_j x_{k,j},$$

where each β_j is a parameter that specifies how the output is related to the j input. Its accuracy is limited, however, by the assumption of a linear relationship between the input factors and the output (gas demand in this case). For the daily demand model, β_0 may represent base load, β_1 may represent the use per heating degree day factor, and $x_{i,1}$ may represent Heating Degree Days, etc. GasDayTM has up to an eight day forecasting horizon, with a separate MLR for each forecast horizon.

2.2 Artificial neural networks Artificial Neural Networks (ANN) [36, 37, 41] are mathematical models which can approximate any (non-linear) continuous function arbitrarily well [23, 24]. The ANN acquires knowledge through a training process [42]. Modelers of gas consumption have been attracted to ANN's because of this capability of mapping unknown nonlinear relationships between inputs and the output [27]. In particular, the nonlinear properties of the ANN allow the direct input of temperature, wind speed, and prior day temperatures into the ANN nodes without accounting for interactions and the nonlinear response of these impacts [8, 9, 10]. In addition, the training process builds an input-output relationship that interpolates well to a situation that may not exactly match the training data.

However, while an ANN is quite good at interpolating a solution that was not presented during training, it is not as good at extrapolating outside the domain of the training knowledge. For example, in the gas estimation problem, this means that if the ANN model was not trained with historical data from days of extreme weather, the model may not perform well on such days. GasDayTM uses a separate ANN for each forecasting horizon.

2.3 Dynamic model adaptation Combining multiple forecasts from models such as artificial neural networks or multiple linear regression can reduce errors arising from faulty assumptions, bias, or mistakes in data [21]. Bates and Granger [6] suggest that combining several forecasts together tends to decrease forecasting error because the combined forecast has equal to or frequently less variance than each of the component forecasts, and Dickinson [14] provides a mathematical proof of this. Armstrong [2] surveyed research on combining forecasts over the last 40

years, concluding that to obtain the best combined forecast accuracy the following guidelines should be considered.

- use different component forecast methods,
- use at least five component forecasts when possible,
- use equal weights unless you have strong evidence to support unequal weighting of forecasts,
- use trimmed means,
- use different data, and
- use the track record and domain knowledge to vary the combination weights.

Natural gas forecasting is an ideal case for combining forecasts, because the forecaster is not always certain which forecasting model is most accurate. As discussed earlier, linear regression extrapolates better than ANNs, but the ANNs often perform better on days similar to ones in its training set. While our two component models (MLR and ANN) each produce estimates of consumption, a weighted combination of these models often yields improved results over the best of the component model estimates. The combination of component model estimates also helps to hedge the forecast since it tends on average to produce estimates that deviate less from the actual consumption.

The model parameters for a ANN are fixed each time the model is retrained. However, combining techniques can allow the forecasting model to adapt dynamically each time it is run, by dynamically updating weights, for example, adjusting the weights to compensate for load growth or behavioral changes in gas consumption between offline retraining of the underlying models. Applying dynamic model adaptation to a daily load forecasting system can both reduce the daily average error and reduce the worst case errors caused by unusual days not observed in the training set.

3 Factors that affect daily demand In this section, we discuss many factors that affect natural gas consumption [11, 30]. Additionally, we show mathematical representations of these factors used as inputs to the GasDayTM MLR and ANN models.

3.1 Modeling temperature effects The most significant factor for modeling natural gas consumption is temperature, since most gas is used for space heating. The daily average temperature and the daily gas consumption for a region in Wisconsin versus day are shown in Figure 1.

(The customer demand has been scaled to protect proprietary information.)




FIGURE 1: Daily average temperature and daily gas consumption for a region in Wisconsin vs. day.

When temperatures are cold, as temperature increases, gas consumption decreases in a nearly linear way, although once the ambient temperature reaches approximately 55 to 65 degrees Fahrenheit, consumption levels off. Once the average temperature reaches a certain temperature, space heating no longer occurs; consumption levels are near some constant value known as base load. This nonlinear characteristic was observed long ago and used to define the *Heating Degree Day* (*HDD*) [3, 4, 13] as

$$HDD_k = \max(0, T_{\text{ref}} - T_k),$$

where T_k is the average temperature for the k^{th} day, and T_{ref} is the reference temperature, historically set to 65°F or 18°C. Gas consumption versus average temperature for individual days has been plotted in Figure 2, which illustrates that gas consumption is approximately proportional to *HDD*.




FIGURE 2: The scatter plot of gas consumption vs. temperature.

The 2-Parameter Model: The demand for the k^{th} day can be estimated as the base load plus a heat load factor times the HDD_k :

$$\hat{S}_k = \beta_0 + \beta_1 HDD_k.$$

The 3-Parameter Model: Over time, the HDD reference temperature has changed. One way to adjust for this change is to add a second HDD factor to the model. Here we add the second HDD factor with a reference temperature of 55°F , which automatically generates an optimal HDD reference temperature as illustrated in Figure 3.

The 5-Parameter Model: Heat loss is a dynamic process. Adding a term for the HDD of the previous day can improve the accuracy of the model. Adding a cooling degree term $CDD_k = \max(T_k - T_{\text{ref}}, 0)$ can also improve the accuracy of the model. Both terms have been added to the model

$$\hat{S}_k = \beta_0 + \beta_1 HDD_k^{65} + \beta_2 HDD_k^{55} + \beta_3 \Delta HDD_k + \beta_4 CDD_k^{65},$$

where $\Delta HDD_k = HDD_k - HDD_{k-1}$.




FIGURE 3: Adding a second *HDD* term has the effect of allowing algorithmic adjustment of the *HDD* reference temperature.

3.2 Modeling wind effects Another important factor is wind, because buildings lose more heat on a windy day than on a calm day. Wind could be added as another term to the models above, but then the wind effect would be the same at all temperatures, while it is well known that the impact of wind increases with *HDD*. A common method that works well is to use *Heating Degree Days adjusted for Wind (HDDW)*. If *WS* is wind speed in mph,

$$HDDW = \begin{cases} \left(\frac{WS + 152}{160} \right) \times HDD, & WS \leq 8 \\ HDD, & WS = 8 \\ \left(\frac{WS + 72}{80} \right) \times HDD, & WS > 8. \end{cases}$$

3.3 Previous day demand Typically, the load forecasts are made for the coming day before the current day's gas day is complete. Thus, the current day's demand is not known. However, yesterday *is* over, so the flow for that day may be known. Adding this and earlier daily flows as inputs to the forecast model, making it autoregressive, can reduce forecast error significantly.

3.4 Modeling day of the week effects Gas consumption varies by the day of the week. For example, on weekends, as residential consumption increases, demand is typically more than offset by decreased consumption of both commercial and industrial consumption. Gas load forecasters have used many techniques to try to capture this effect, primarily by adding day-of-the-week indicator model inputs.

Weekday/Weekend indicator: A binary indicator variable can be added to the model to distinguish weekdays from weekends. That is, the variable Weekend is 1 on Saturdays and Sundays and 0 on the other days of the week. This term can be added to any of the models described above.

Friday indicator: Since the industry definition of a gas day for a Friday includes the Saturday morning start-up, typically demand for gas day Friday is lower than the other weekdays, yet higher than Saturday and Sunday demands. This effect varies from region to region at an LDC and certainly across the country. This effect can be accounted for by setting the indicator variable to a number between 0 and 1 on Fridays.

Sine/cosine indicators: Periodic phenomena can be represented by Fourier series [38]. The days of the week are periodic with a period of seven days, so we can use a day-of-the-week DOW variable to represent the fundamental seven day frequency, 1 = Sunday, 2 = Monday, etc.:

$$\sin\left(\frac{2\pi DOW}{7}\right), \quad \cos\left(\frac{2\pi DOW}{7}\right).$$

Seven day lag: Another technique for improving the demand forecast includes both the demand and the temperature/ HDD for the day seven days ago, unless the day seven days earlier was a holiday.

3.5 Holidays and days around holidays Holidays and days near holidays typically have lower demands than if the day was not a holiday. One approach that can be used with or without the above mentioned day of week adjustment is to average the residual errors in the training data on specific holidays and adjust the demand forecast. For example, if, after parameterizing our model, we evaluate the model on all of the New Years Days in the training set and calculate the forecast error as the demand forecast minus the actual flow, and calculate the average errors, we can subtract this average error to adjust the forecast on New Years Day.

Holiday adjustments: Another common way to predict gas consumption for a holiday is to pretend that the day is a Saturday. Days near holidays also can be adjusted, i.e., the day before a holiday can be set to a Friday, or when a holiday falls on a Monday, set the Sunday to Saturday and set the holiday (Monday) to a Sunday. Models that use demands from previous days are biased low on days after holidays, as the low holiday demand is now an input to the model. This low holiday demand can be adjusted by adding in the average error to make it act like a non-holiday.

3.6 Other factors Many other potential factors exist, such as solar radiation, cloud cover, precipitation, dew point, direction of the wind, tap water temperature, bill shock, occupancy rates, industrial production rates, and other econometric factors, to name a few. Some of these factors can be measured directly, while others cannot, or at least, cannot be measured reasonably. Solar radiation and cloud cover affect temperatures throughout the day. For example, the evening temperature decrease is less on a cloudy day than on a clear day. Wind direction has an effect, especially in coastal regions next to oceans or the Great Lakes. The dew point is a measure of humidity, and more gas tends to be consumed on humid days. Precipitation measures rain and snow fall. Industrial production and other economic factors affect gas consumption especially in parts of the country where there is a large industrial concentration. For example, in Michigan the economic recession caused a decrease in auto production, reducing gas consumption. Bill shocks, as experienced in 2005 when hurricanes hit the Gulf of Mexico, cause people to turn down their thermostats in an effort to reduce their bills after having paid a large gas bills for their prior billing period or in response to media coverage.

4 Data quality When building models from historical data, the data quality of the training data set is critical. Most model fitting algorithms, including MLR and the ANN training methods discussed above, are designed to minimize the error standard deviation, a form of squared error. If the training data contains errors, the model does not fit well.

Data cleaning: The best scenario is to start with good data, but enough good historical data is not always available. We can use preliminary demand forecast models to detect anomalous data, which can be confirmed, and corrected or discarded before final models are developed [26, 28, 39, 40].

Data disaggregation: Similarly, for some LDC's, the only historical demand data available may be (approximately) monthly, although daily demand forecasts are required. It often is possible to use preliminary demand forecast models to disaggregate monthly data into approximate historical daily data before final models are developed [45]. Similar techniques can be used to give good hourly demand forecasts in the presence of unreliable hourly flow data [35].

Number of training days: Insufficient historical data can introduce problems in training models to predict gas consumption. When training ANN's, heuristically, ten times as many training set vector pairs as weights in the ANN are needed. Otherwise, the ANN will "memorize" the training vector pairs and will not generalize the trends in the data well. This memorizing phenomena is known as over-training or over-fitting. Similar problems occur with linear regression models if the training data set is not large enough or sufficiently rich.

Growth in the customer base: Models are *developed* from a very large experience base, but models for an LDC are trained on historical data from that LDC. The most recent years of gas consumption history tend to be most relevant. The older data is not a good indication of the current customer base characteristics due to both customer base growth ("growth" can be negative) and demand-side management. Because of this non-stationary customer base, building a model to predict demand for the next heating season is difficult, but growth adjustments can be calculated [19, 20, 22, 25, 44].

Let us consider an example to illustrate this. Suppose a model is built using data from the most recent five years from an operating area with substantial growth. If all days in the training data set are equally weighted, the model best predicts the load for the average customer base in the training data set. The residual errors of the model is smallest for the middle year. The errors tend to be positive (larger predictions than actual demand) over the first two years of the training data, and tend to be negative (smaller predictions than actual demand) over the last two years of the training data. Our goal is to build a model to predict demand for the *coming* heating season, but the model best predicts the heating season three years prior.

This problem can be partially overcome by "growing" the older historical data [12]. A simple way to grow historical data is to make it all look like it occurred during the most recent heating season. This is accomplished by calculating linear regression models for each heating

season. The demands from heating seasons before the most recent can be adjusted by adding a base-load factor to each day to make the base loads the same as the most recent season and by adding additional demand proportional to the *HDD* to each day to make the use per *HDD* factor the same as the use per *HDD* for the most recent season. For example, using only 2008–2009 data, we built a 2-parameter model

$$\hat{S}_k = \beta_0^{2009} + \beta_1^{2009} HDD_k,$$

and using only 2007–2008 data, we built a 2-parameter model

$$\hat{S}_k = \beta_0^{2008} + \beta_1^{2008} HDD_k.$$

We then grew the 2007–2008 heating season data as

$$S_k^{\text{new}} = S_k + (\beta_0^{2009} - \beta_0^{2008}) + (\beta_1^{2009} - \beta_1^{2008}) HDD_k.$$

The “new” 2007–2008 demand data has the same base load and heat load factor as the 2008–2009 data, but it is appropriate for the weather of 2007–2008.

Flow vs. demand: Models built using flow (consumption) data predict flow. Models built using demand data predict demand. On most days, flow equals demand. However, on days when the LDC interrupts customers, or injects from storage, or when a hurricane disrupts an entire region, the gas that flows through the city gate stations is less than the demand for gas.

Figure 4 shows the actual flow versus temperature for a different operating area than Figure 2, which contains many interruptible customers. The “bend over” effect at colder temperatures is caused by customers being interrupted. To make a flow forecasting model predict demand, the historical training data must be augmented with estimated interrupted flow, so that the model built using this data predicts demand.

Operating areas: Forecast accuracies often can be improved by subdividing the region for which demand forecasts are required into smaller operating areas and forecasting each area with separately trained models. Smaller areas may benefit from more accurate average weather forecasts, from a more homogenous customer base, or from other factors.

Multiple weather stations: Forecast accuracies often can be improved by using carefully tuned weighted averages of weather forecasts from multiple stations in or near the target operating area.




FIGURE 4: Flow versus temperature for an operating area with interruptible customers.

Generalization, interpolation, and extrapolation: Claims have been made that ANN's are excellent generalizers [24, 47]; that an ANN can learn general trends from a training data set and then can make valid estimates for an input that it has not seen before [7, 43]. This is true if the input is similar to inputs in the training data, but it is false if the input is not close to any of the inputs in the training data. A better way to state the capabilities of an ANN is that it interpolates well, but in general, it extrapolates unpredictably [5, 16]. In contrast, the linear regression model extrapolates very predictably [31], and in the gas demand forecasting case, quite well.

This implies that the ANN model forecasts gas demand estimates well on days that are similar to historical days in the training set, and not as well on days that are not similar to those in the training set. This rightly brings up concerns for demand estimations on peak days and even uncommon days (days that are significantly colder, warmer or windier than normal, days that are much warmer or colder than the previous day, etc.).

Figure 5 shows temperature versus wind for the training set (12 Nov. 1994 to 31 May 1997) and the testing set (1 June 1997 to 31 May 1998) for an ANN trained for the 1997–1998 heating season. Even though the 1997–1998 heating season was mild (El-Niño), the testing set (the

heating season) contained many windy days that were not similar to any days in the training set, which happened to include only two windy days. Three approaches to solve this problem are (1) use more (and older) training data, (2) fabricate additional training data, and (3) use surrogate data from another operating area with similar customer base and temperature characteristics.




FIGURE 5: Temperature versus wind for the training and testing data sets.

Data weighting: In Milwaukee, WI, today's average temperature is within 8°F of yesterday's temperature about 80% of the time. However, the challenge in natural gas demand forecasting is not the typical day, but the *unusual* day. If we build models using equal weighting of all the data, days where today's temperature is more than 8°F different yesterday's temperature will be weighted one fourth as important as the other days. However, these are the days that LDC's need good forecasts, and forecasting is most difficult. When modeling natural gas, a gas forecaster should consider identifying unusual days, such as today much colder or warmer than yesterday, much colder or warmer than normal, much windier than normal, etc., so these days can be weighted more heavily in the model training process.

5 GasDay performance The following presentation of GasDayTM forecasting results is performed on 14 different operating areas for a utility in the US. The models that generated these forecasts were trained on September 14, 2009 so results shown from October 2009 through July 2010 are exclusive of the training data set. Their flow has been scaled to protect their proprietary data. Figure 6 shows the scaled flow and the GasDayTM estimated flow for a one-day-ahead forecast for January 2010.

Figure 7 shows the temperature and one-day-ahead temperature forecasts for January, which corresponds to the flow and flow estimate in Figure 6. Figures 6 and 7 show time series for operating area one. For January 7th, 16th, 19th, and 20th, GasDayTM had large forecast errors, as shown in in Figure 6. On these days, the weather forecast error was also several degrees. Hence, the accuracy of GasDay'sTM forecasts are dependent on the accuracy of the weather forecasts used.

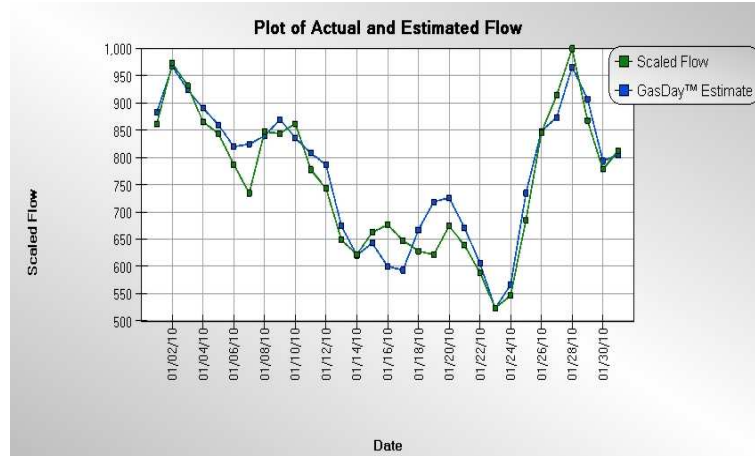


FIGURE 6: Flow and GasDayTM one-day-ahead estimated flow

Tables 1 and 2 show Mean Absolute Percent Error (MAPE) by month for the 14 different operating areas. MAPE for the period of October 2009 through July 2010 is also reported. Table 1 and 2 show results for the Linear Regression (LR), Artificial Neural Network (ANN), and GasDay'sTM combined estimate (GD). We empirically observe the same conclusions Bates and Granger [6] and Dickinson [14] assert. For the




FIGURE 7: Temperature and one-day-ahead temperature forecast

period of October 2009 through July 2010 on 11 of 14 operating areas, GasDay'sTM estimate is better than the best of the two component models, and all 14 GasDay estimates are better than the worst component model. Operating areas 3, 9, and 14 contain large concentrations of industrial customers that have less heating load sensitivity, making them harder to forecast.

6 Summary In this paper, we have emphasized the importance for LDCs to make accurate natural gas demand forecasts and the financial consequences to their customers if they do not. Additionally, we described two important model fitting algorithms used by GasDayTM to forecast daily natural gas demand: multiple linear regression and artificial neural networks. The impacts of temperature, wind, prior day weather, previous day demands, day of the week, and holidays on gas consumption have been discussed, along with common data quality issues such as the length of the training data set, the differences between flow and demand, and customer base growth are also discussed. In addition, we described the models and variables that are used by GasDayTM to forecast LDCs consumption and the data quality issues that must be addressed before good models can be trained. A survey of GasDayTM performance results show that GasDayTM forecasts gas consumption well.

When applied properly, GasDayTM using a combination of multiple linear regression and artificial neural networks is a very accurate tool for forecasting daily gas demand.

	1			2			3			4		
Month	LR	ANN	GD	LR	ANN	GD	LR	ANN	GD	LR	ANN	GD
Oct.'09	8.83	9.36	8.72	12.40	11.53	10.02	14.59	16.61	13.30	6.52	5.76	6.06
Nov.'09	5.28	5.52	5.52	10.33	8.79	9.57	17.58	15.26	15.18	6.24	6.79	6.20
Dec.'09	4.48	5.21	5.05	5.45	5.09	5.72	11.62	14.33	11.54	3.68	3.89	4.13
Jan.'10	4.80	4.73	4.28	5.38	4.92	4.73	10.94	13.75	12.71	4.06	3.87	4.01
Feb.'10	7.57	6.47	4.76	3.22	3.22	3.86	9.79	7.61	8.68	4.90	3.29	3.86
Mar.'10	6.30	6.72	6.68	7.29	7.41	6.48	14.31	13.20	12.90	4.97	5.01	5.12
Apr.'10	9.86	10.27	8.66	10.64	8.65	11.41	29.84	31.99	29.34	8.14	8.08	8.66
May'10	5.81	6.23	6.07	10.84	14.79	12.16	31.12	28.93	27.43	6.94	5.51	6.10
Jun.'10	5.82	5.35	6.44	15.77	15.53	10.70	40.26	31.46	32.73	9.29	8.80	8.66
Jul.'10	5.29	5.36	6.27	20.23	20.81	11.99	32.91	31.13	31.67	5.82	5.96	5.46
Oct.'09 to Jul.'10	6.39	6.52	6.25	10.20	10.13	8.69	21.33	20.50	19.59	6.05	5.70	5.83
	5			6			7			8		
Month	LR	ANN	GD	LR	ANN	GD	LR	ANN	GD	LR	ANN	GD
Oct.'09	3.16	5.06	3.67	9.82	8.26	8.86	12.54	10.04	9.25	5.58	5.87	6.23
Nov.'09	8.66	13.12	9.73	5.30	5.06	4.42	11.21	11.98	8.86	11.66	9.65	9.78
Dec.'09	6.28	5.33	6.13	5.89	6.12	5.83	6.47	10.25	5.66	5.40	5.41	5.84
Jan.'10	5.78	6.51	6.15	4.60	4.34	4.41	4.79	4.81	4.01	3.42	2.35	2.70
Feb.'10	2.48	3.56	2.95	5.60	5.63	5.25	4.06	3.95	4.07	4.81	4.68	4.83
Mar.'10	4.18	4.40	3.02	7.27	7.12	7.08	8.51	12.12	8.22	4.14	4.16	4.62
Apr.'10	5.66	5.72	5.34	6.04	6.09	6.19	9.42	12.95	7.38	6.53	6.77	6.71
May'10	5.36	6.45	5.81	8.26	7.45	7.64	12.92	13.28	13.73	6.37	6.42	6.20
Jun.'10	3.16	3.33	3.44	7.30	8.02	7.38	11.19	7.19	8.30	7.66	7.96	6.23
Jul.'10	10.78	9.79	9.97	14.92	14.21	15.36	11.75	9.86	10.47	7.61	8.10	4.29
Oct.'09 to Jul.'10	5.58	6.34	5.64	7.53	7.25	7.27	9.33	9.69	8.03	6.31	6.13	5.73

TABLE 1: MAPE performance evaluation by month for the 2009/2010 heating season for 14 operating areas.

	9			10			11		
Month	LR	ANN	GD	LR	ANN	GD	LR	ANN	GD
Oct.'09	8.26	7.04	5.82	9.47	8.35	7.03	14.40	14.87	13.90
Nov.'09	7.02	7.57	6.39	8.11	10.36	9.51	8.41	9.96	8.09
Dec.'09	5.26	4.61	5.16	6.77	10.34	7.04	5.00	5.21	5.89
Jan.'10	5.33	4.99	5.17	4.86	4.89	5.70	4.48	4.25	4.32
Feb.'10	7.93	6.32	6.21	4.53	3.77	4.31	6.73	5.69	4.97
Mar.'10	11.66	9.61	8.68	8.70	8.84	7.80	5.78	5.88	6.45
Apr.'10	15.33	10.48	9.44	14.01	18.08	9.89	10.41	9.95	10.12
May'10	29.63	23.11	19.27	12.98	11.39	6.40	6.01	6.05	6.85
Jun.'10	28.87	26.20	17.65	13.83	21.50	8.54	6.55	5.72	6.61
Jul.'10	21.91	23.31	13.01	14.48	14.25	13.34	4.81	4.82	4.89
Oct.'09 to Jul.'10	14.15	12.36	9.70	9.80	11.20	7.98	7.25	7.24	7.22
	12			13			14		
Month	LR	ANN	GD	LR	ANN	GD	LR	ANN	GD
Oct.'09	10.67	9.18	8.09	9.68	8.99	9.65	23.26	25.44	16.14
Nov.'09	6.33	5.28	5.58	11.68	10.73	9.19	16.30	17.44	16.70
Dec.'09	4.17	3.56	3.48	4.63	4.61	4.86	18.08	19.91	19.71
Jan.'10	4.59	4.38	4.02	3.92	3.76	4.17	9.49	10.14	8.28
Feb.'10	3.42	3.39	4.22	5.99	6.70	6.20	9.28	10.02	8.29
Mar.'10	6.39	6.59	5.49	5.67	5.77	5.99	10.98	10.45	9.66
Apr.'10	8.43	9.38	7.24	10.07	7.81	8.54	25.01	19.66	19.19
May'10	7.48	8.53	3.99	5.68	5.67	5.53	18.99	14.47	14.95
Jun.'10	8.13	7.20	4.92	4.56	4.67	4.59	17.20	15.26	16.04
Jul.'10	10.77	8.30	5.06	4.95	4.79	4.65	19.22	14.00	15.42
Oct.'09 to Jul.'10	7.07	6.60	5.21	6.67	6.33	6.33	16.83	16.00	14.47

TABLE 2: MAPE performance evaluation by month for the 2009/2010 heating season for 14 operating areas.

REFERENCES

1. American Gas Association, (2010), <http://www.aga.org>.
2. J. S. Armstrong, *Principles of Forecasting a Handbook for Researchers and Practitioners*, Kluwer Academic Publishers, Boston, 2001.
3. J. G. Asbury, C. Maslowski and R. O. Mueller, *Solar availability for winter space heating: an analysis of SOLMET data, 1953 to 1975*, Science **206**(4419) (1979), 679–681.

4. D. G. Baker, *Effect of observation time on mean temperature estimation*, J. Appl. Meteor. **14**(4) (1975), 471–476.
5. E. Barnard and L. F. A. Wessels, *Extrapolation and interpolation in neural network classifiers*, IEEE Control Syst. Mag. **12**(5) (1992), 50–53.
6. J. M. Bates and C. W. J. Granger, *The combination of forecasts*, Oper. Res. Q. **20** (1969), 451–468.
7. D. S. Broomhead and D. Lowe, *Multivariable functional interpolation and adaptive networks*, Complex Syst. **2**(3) (1988), 321–355.
8. R. H. Brown, I. Matin and X. Feng, *Development of feed-forward neural network models for gas short-term load forecasting*, 222–229, Proc. Adaptive Control Systems Technology Symposium; Theme: High Tech Controls for Energy and Environment, Pittsburgh, PA, 1994.
9. R. H. Brown, I. Matin, P. Kharouf and L. P. Piessens, *Development of artificial neural network models to predict daily gas consumption*, AGA Forecasting Rev. **5** (1996), 1–22.
10. R. H. Brown, T. M. Richardson and J. E. Buchanan, *Forecasting daily sendout demand with artificial neural networks*, 131–139, Amer. Gas Assoc. Operating Section Pre-Print Proc., Cleveland, OH, 1999.
11. R. H. Brown, *OK, who is playing with the thermostat?*, Gas Forecaster's Forum, Albuquerque, NM, October 2003.
12. R. H. Brown, Y. Li, B. Pang, S. R. Vitullo and G. F. Corliss, *Detrending daily natural gas demand data using domain knowledge*, 10, Proc. 30th Int. Symposium on Forecasting, San Diego, CA, 2010.
13. P. M. Dare, *A study of the severity of the midwestern winters of 1977 and 1978 using heating degree days determined from both measured and wind chill temperatures*, Bull. Amer. Meteor. Soc. **62**(7) (1981), 974–982.
14. J. P. Dickinson, *Some statistical results in the combination of forecasts*, Oper. Res. Q. **24**(2) (1998).
15. N. R. Draper and H. Smith, *Applied Regression Analysis*, 3rd ed., John Wiley & Sons, New York, 1998.
16. W. Duch and G. H. F. Dierksen, *Neural networks as tools to solve problems in physics and chemistry*, Comput. Phys. Comm. **82**(2–3) (1994), 91–103.
17. Energy and Environmental Analysis Inc., *The Impacts of Recent Hurricanes on U.S. Gas Markets for the Upcoming Winter*, Technical Report, 2005.
18. A. S. Goldberger, *Topics in Regression Analysis*, Macmillan, New York, 1968.
19. T. Haida and S. Muto, *Regression based peak load forecasting using a transformation technique*, IEEE Trans. Power Syst. **9**(4) (1994), 1788–1794.
20. T. Haida, S. Muto, Y. Takahashi and Y. Ishi, *Peak load forecasting using multiple-year data with trend data processing techniques*, Electr. Eng. Japan **124** (1998), 7–16.
21. S. Halim, *Selection of Inputs for Generating Combinatorial Daily Natural Gas Demand Forecasts*, M.S. thesis Marquette University Department of Electrical and Computer Engineering, Milwaukee, WI, 2004.
22. M. He, *Annual Transformation Algorithm for Hourly Data in Natural Gas Consumption*, M.S. thesis, Marquette University Department of Electrical and Computer Engineering, Milwaukee, WI, 2007.
23. J. J. Hopfield, *Neural networks and physical systems with emergent collective computational abilities*, Proc. Nat. Acad. Sci. USA **79** (1982), 2554–2558.
24. K. Hornik, M. Stinchcombe and H. White, *Multilayer feedforward networks are universal approximators*, Neural Networks **2** (1989), 359–366.
25. D. Hughes and R. H. Brown, *Analysis of customer base growth over time*, Gas Forecaster's Forum, Greensboro, NC, 2002.
26. R. O. Kennedy, *Detecting Outliers and Meter Anomalies in Natural Gas Customer Flow Data*, M.S. thesis, Marquette University Department of Electrical and Computer Engineering, Milwaukee, WI, 2006.

27. A. Khotanzad, H. Elragal and T. L. Lu, *Combination of artificial neural network forecasters for prediction of natural gas consumption*, IEEE Trans. Neural Networks **11** (2000), 464–473.
28. S. Kiware, *Detection of Outliers in Time Series Data*, M.S. thesis, Marquette University Department of Mathematics, Statistics, and Computer Science, Milwaukee, WI, 2010.
29. R. R. Levary and B. V. Dean, *A natural gas model under uncertainty in demand*, Oper. Res. **28**(6) (1980), 1360–1374.
30. H. L. Lim and R. H. Brown, *Hourly gas load forecasting model input factor identification using a genetic algorithm*, 670–673, Proc. 44th IEEE Midwest Symposium on Circuits and Systems, Dayton, OH, 2001.
31. H. Lohninger, *Evaluation of neural networks based on radial basis functions and their application to the prediction of boiling points from structural parameters*, J. Chem. Information Comput. Sci. **33**(5) (1993), 736–744.
32. F. K. Lyness, *Gas demand forecasting*, The Statistician **33**(1) (1984).
33. F. K. Lyness, *Consistent forecasting of severe winter gas demand*, J. Oper. Res. Soc. **32**(5) (1981), 347–459.
34. B. M. Marx, *Forecasting Daily Processing Tomato Harvest Tonnage in California*, M.S. thesis, Marquette University Department of Electrical and Computer Engineering, Milwaukee, WI, 2004.
35. B. Marx, *Fitting a Continuous Profile to Hourly Natural Gas Flow Data*, Ph.D. thesis, Marquette University Department of Electrical and Computer Engineering, Milwaukee, WI, 2007.
36. W. S. McCulloch and W. Pitts, *A logical calculus of the ideas immanent in nervous activity*, Bull. Math. Biophys. **5** (1943), 115–133.
37. M. L. Minsky and S. A. Papert, *Perceptrons*, The MIT Press, Cambridge, MA, 1969. (Expanded ed. 1990.)
38. A. V. Oppenheim, A. S. Willsky and S. H. Nawab, *Signals and Systems* 2nd ed., Prentice Hall, Upper Saddle, NJ, 1997.
39. Ronald K. Pearson, *Outliers in process modeling and identification*, IEEE Trans. Control Syst. Technol. **10**(1) (2002), 55–63.
40. Ronald K. Pearson, *Mining Imperfect Data*, SIAM, U.S.A., 2005.
41. F. Rosenblatt, *Principles of Neurodynamics: Perceptrons and the Theory of Brain Mechanisms*, Spartan, New York, 1962.
42. D. E. Rumelhart, G. E. Hinton, and R. J. Williams, *Learning internal representations by error propagation*, Parallel Distributed Processing, D. E. Rumelhart, J. L. McClelland, and the PDP Research Group, The MIT Press, Cambridge, MA, 1986, 318–362.
43. D. Suter, *Neural net surface interpolation*, Proceedings of the 1987 International Conference on Systems, Man, and Cybernetics, IEEE **1** (1987), 118–123, New York, NY., USA.
44. A. Taware, *Forecasting and Identification Methods Applied to Gas Load Estimation Problems*, M.S. thesis, Marquette University Department of Electrical and Computer Engineering, Milwaukee, WI, 1998.
45. S. R. Vitullo, *Disaggregating Interval Time-Series Data Applied to Natural Gas Flow Estimation*, M.S. thesis, Marquette University Department of Electrical and Computer Engineering, Milwaukee, WI, 2007.
46. T. H. Welch, J. G. Smith, J. P. Rix and R. D. Reader, *Meeting seasonal peak demand for natural gas*, Operational Research Quarterly **22** (1971), 93–106.
47. A. Weiland and R. Leighton, *Geometric analysis of neural network capabilities*, 385–392, IEEE First International Conference on Neural Networks, SOS Printing 3, San Diego, CA, 1987.

DEPARTMENT OF ELECTRICAL AND COMPUTER ENGINEERING
P. O. Box 1881, MARQUETTE UNIVERSITY, MILWAUKEE, WI 53201-1881

E-mail address: steven.vitullo@marquette.edu

E-mail address: ronald.brown@marquette.edu

E-mail address: george.corliss@marquette.edu

E-mail address: brian.marx@marquette.edu

